

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.model_selection import train_test_split as tts
from sklearn.metrics import r2_score as r2, mean_squared_error as mse
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df = pd.read_csv('Walmart_Store_sales.csv')
df.head()
```

```
Out[2]:
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106

```
In [3]: df = df.rename(columns={'Store':'store', 'Date':'date', 'Weekly_Sales':'weekly_sales', 'Holiday_Flag':'holiday_flag',
'Fuel_Price':'fuel_price', 'CPI':'cpi', 'Unemployment':'unemployment'})
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   store           6435 non-null   int64
1   date            6435 non-null   object
2   weekly_sales    6435 non-null   float64
3   holiday_flag    6435 non-null   int64
4   temp            6435 non-null   float64
5   fuel_price      6435 non-null   float64
6   cpi             6435 non-null   float64
7   unemployment    6435 non-null   float64
dtypes: float64(5), int64(2), object(1)
memory usage: 402.3+ KB
```

1 Basic Statistics Task

1.1 Store with Maximum Sale

```
In [5]: print('Store with Highest Sale')
print((df.groupby('store').sum('weekly_sales').sort_values('weekly_sales',ascending=False)

Store with Highest Sale
store
20    3.013978e+08
4     2.995440e+08
14    2.889999e+08
Name: weekly_sales, dtype: float64
```

1.2 Highest Standard Deviation in weekly_sales for top 3 Stores

```
In [6]: std_df= df.groupby('store').agg({'weekly_sales':['std','mean']}).sort_values(('weekly_
std_df['coefficient_of_variance']= std_df[['weekly_sales','std']]/std_df[['weekly_sale
print('\nStore with Highest Standard Deviation in Sales')
print(std_df.iloc[:3])

Store with Highest Standard Deviation in Sales
              weekly_sales              coefficient_of_variance
              std              mean
store
14    317569.949476  2.020978e+06              15.713674
10    302262.062504  1.899425e+06              15.913349
20    275900.562742  2.107677e+06              13.090269
```

1.3 Highest Growth rate for top 3 Stores in Q3 2012

```
In [7]: df['date']= pd.to_datetime(df['date'],format='%d-%m-%Y')
```

```
In [8]: from datetime import datetime as dt
q2_start= dt(2012,4,1)
q2_end= dt(2012,6,30)
q3_start= dt(2012,7,1)
q3_end= dt(2012,9,30)
```

```
In [9]: q2_df= df[df['date'].between(q2_start,q2_end)][['store','weekly_sales']]
q3_df= df[df['date'].between(q3_start,q3_end)][['store','weekly_sales']]
```

```
In [10]: quarters_df= q3_df.groupby('store').sum('weekly_sales').sort_index()
quarters_df= quarters_df.rename(columns={'weekly_sales':'q3_sales'})
quarters_df['q2_sales']= q2_df.groupby('store').sum('weekly_sales')['weekly_sales'].sort_index()
quarters_df['growth_rate']= (quarters_df.q3_sales-quarters_df.q2_sales)/quarters_df.q2_sales
quarters_df.sort_values('growth_rate',ascending=False).iloc[:3]
```

```
Out[10]:
```

	q3_sales	q2_sales	growth_rate
store			
7	8262787.39	7290859.27	0.133308
16	7121541.64	6564335.98	0.084884
35	11322421.12	10838313.00	0.044666

1.4 Holiday and Non-Holidays AVG Sales comparison

```
In [11]: df.groupby('holiday_flag')['weekly_sales'].mean()
```

```
Out[11]: holiday_flag
0      1.041256e+06
1      1.122888e+06
Name: weekly_sales, dtype: float64
```

```
In [12]: ch1 = dt(2010,12,31)
ch2 = dt(2011,12,30)
ch3 = dt(2012,12,28)
ch4 = dt(2013,12,27)

th1= dt(2010,11,26)
th2= dt(2011,11,25)
th3= dt(2012,11,23)
th4= dt(2013,11,29)

la1= dt(2010,9,10)
la2= dt(2011,9,9)
la3= dt(2012,9,7)
la4= dt(2013,9,6)

su1= dt(2010,2,12)
su2= dt(2011,2,11)
su3= dt(2012,2,10)
su4= dt(2013,2,8)
```


```
In [13]: christmas_mean_df= df[(df['date']== ch1) | (df['date']== ch2) | (df['date']== ch3) | (
thanksgiving_mean_df= df[(df['date']== th1) | (df['date']== th2) | (df['date']== th3)
labourday_mean_df= df[(df['date']== la1) | (df['date']== la2) | (df['date']== la3) | (
superbowl_mean_df= df[(df['date']== su1) | (df['date']== su2) | (df['date']== su3) | (
```

```
In [14]: dict_mean_sales = {'Thanksgiving_Mean_Sales': thanksgiving_mean_df['weekly_sales'].mean(),
                          'SuperBowl_Mean_Sales': superbowl_mean_df['weekly_sales'].mean(),
                          'LabourDay_Mean_Sales': labourday_mean_df['weekly_sales'].mean(),
                          'Non_Holiday_Weekly_Sales' : df[df['holiday_flag'] == 0]['weekly_sales'].mean(),
                          'Christmas_Mean_Sales' : christmas_mean_df['weekly_sales'].mean()}
```

Only Christmas in Holidays has bad AVG sales than Non-Holiday days

```
In [15]: plt.figure(figsize=(14,8),dpi=100)
plt.bar(x=dict_mean_sales.keys(), height=dict_mean_sales.values())
```

Out[15]: <BarContainer object of 5 artists>



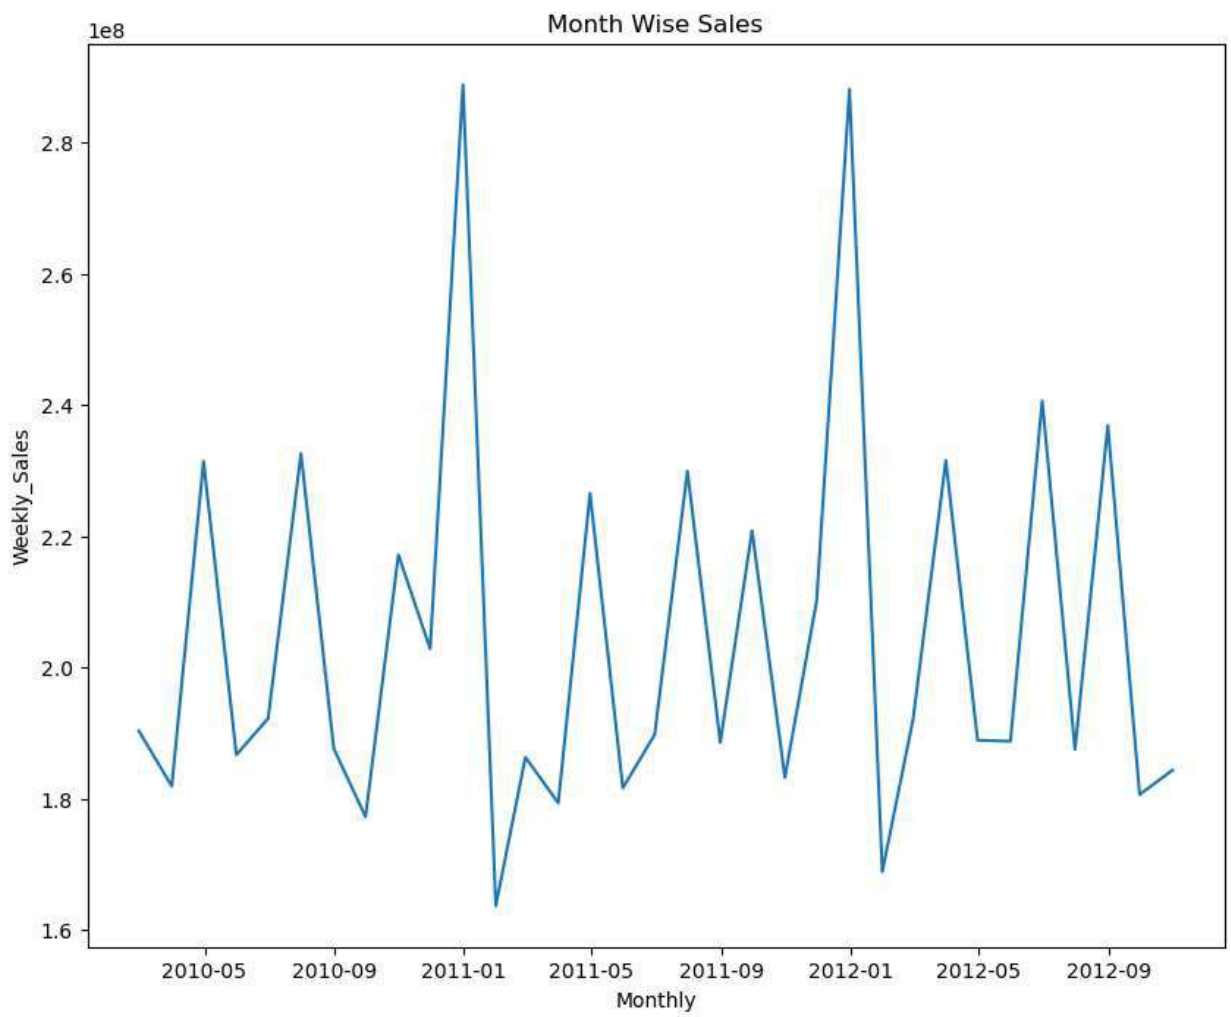
1.5 Monthly, Quarterly, Semester-wise Sale Analysis

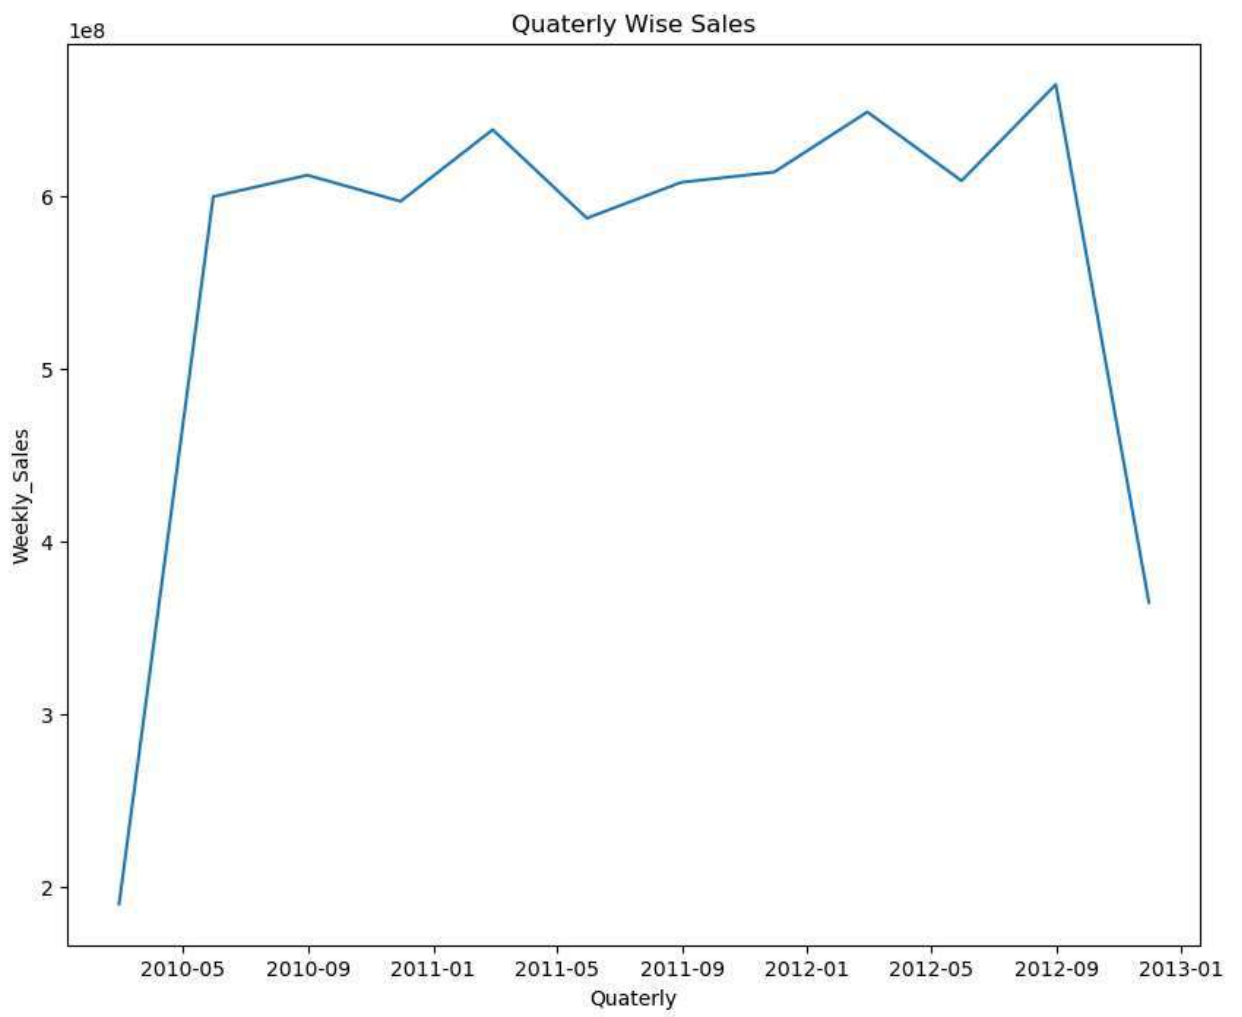
```
In [16]: monthly= df.groupby(pd.Grouper(key='date', freq='1M')).sum()
monthly= monthly.reset_index()
fig,ax= plt.subplots(figsize=(10,8))
plt.plot(monthly['date'],monthly['weekly_sales'])
plt.title('Month Wise Sales')
plt.xlabel('Monthly')
plt.ylabel('Weekly_Sales')

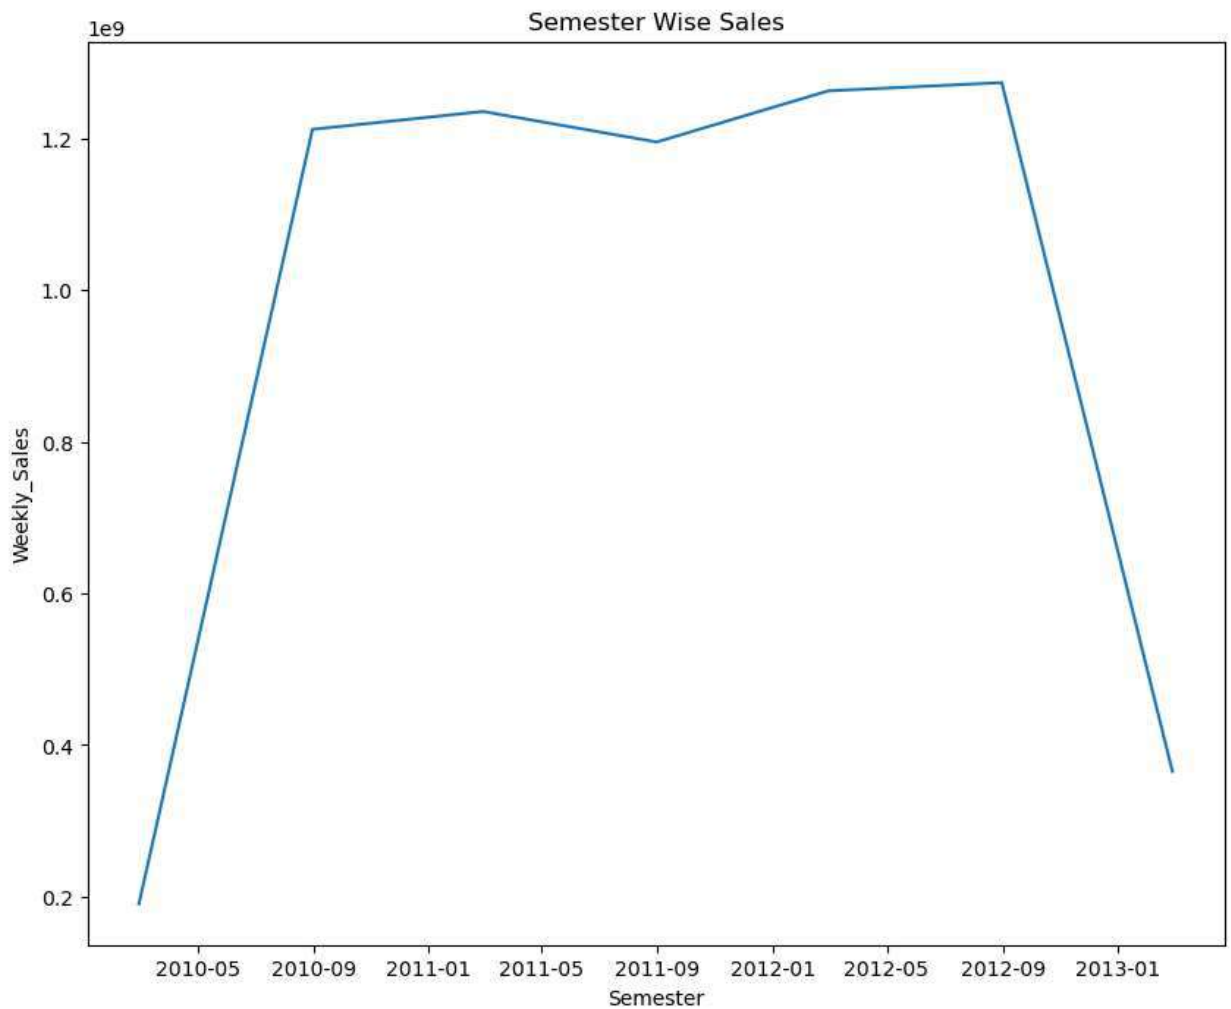
quarterly= df.groupby(pd.Grouper(key='date', freq='3M')).sum()
quarterly= quarterly.reset_index()
fig,ax= plt.subplots(figsize=(10,8))
plt.plot(quarterly['date'],quarterly['weekly_sales'])
plt.title('Quarterly Wise Sales')
plt.xlabel('Quarterly')
plt.ylabel('Weekly_Sales')

semester= df.groupby(pd.Grouper(key='date', freq='6M')).sum()
semester= semester.reset_index()
fig,ax= plt.subplots(figsize=(10,8))
plt.plot(semester['date'],semester['weekly_sales'])
plt.title('Semester Wise Sales')
plt.xlabel('Semester')
plt.ylabel('Weekly_Sales')
```

Out[16]: Text(0, 0.5, 'Weekly_Sales')







2 Statistical Models

2.1 Linear Regression Models

```
In [17]: data= df[df.store==1]
data['date']= np.arange(1,144)
scaler= StandardScaler()
scaler.fit(data)
data.head()
```

C:\Users\admin\AppData\Local\Temp\ipykernel_10744\2008928399.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['date']= np.arange(1,144)

```
Out[17]:
```

	store	date	weekly_sales	holiday_flag	temp	fuel_price	cpi	unemployment
0	1	1	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	2	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	3	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	4	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	5	1554806.68	0	46.50	2.625	211.350143	8.106

```
In [18]: y = data.weekly_sales
X = data.drop(columns=['weekly_sales', 'store', 'date'])
X_train,X_test,y_train,y_test= tts(X,y,test_size=0.2,random_state=42)
```

```
In [19]: model= LinearRegression()
model.fit(X_train,y_train)
y_test_pred= model.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, y_test_pred)))
print('R-Squared:\n',r2(y_test,y_test_pred))
plt.title('Linear regression')
sns.regplot(x=y_test,y=y_test_pred,color='green')
```

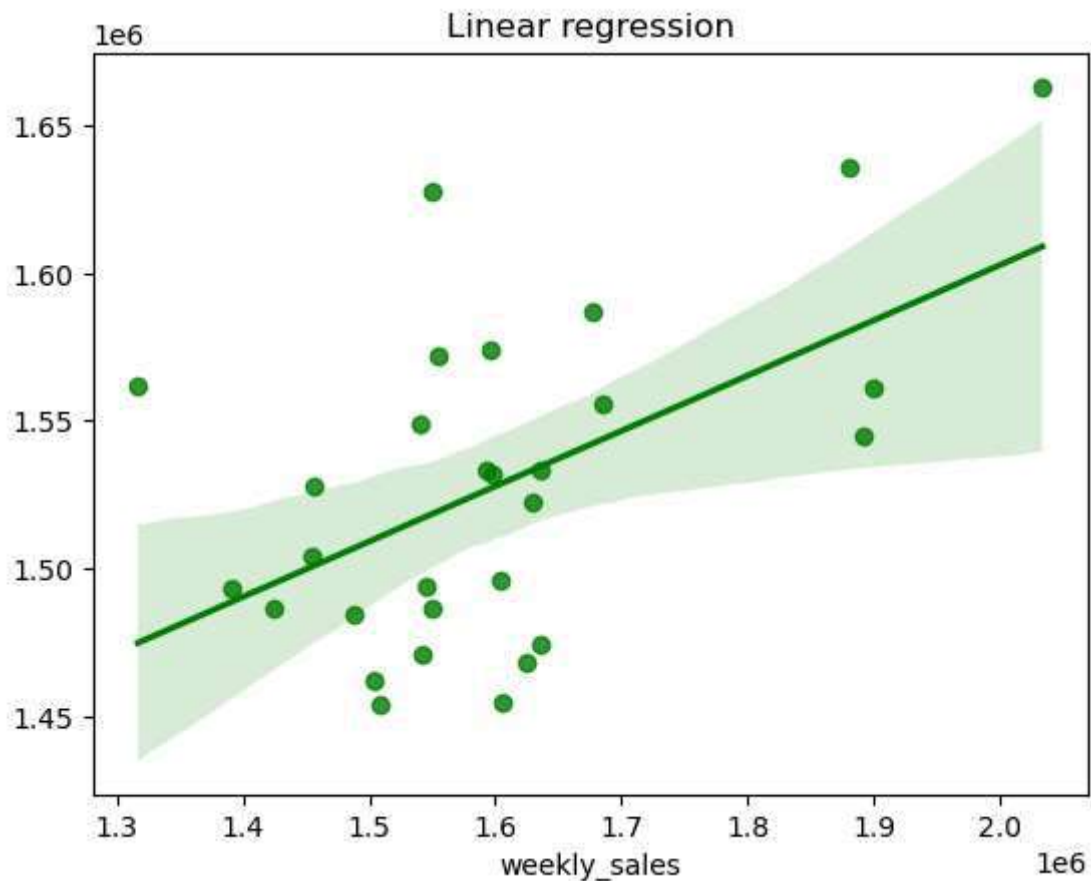
Root Mean Squared Error(RMSE):

152799.6222036366

R-Squared:

0.034707216167411015

```
Out[19]: <Axes: title={'center': 'Linear regression'}, xlabel='weekly_sales'>
```




```
In [20]: lasso_reg=Lasso(alpha=1)
lasso_reg.fit(X_train,y_train)
lasso_y_pred=lasso_reg.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, lasso_y_pred)))
print('R-Squared:\n',r2(y_test,lasso_y_pred))
plt.title('Lasso regression')
sns.regplot(x=y_test,y=lasso_y_pred,color='y')
```

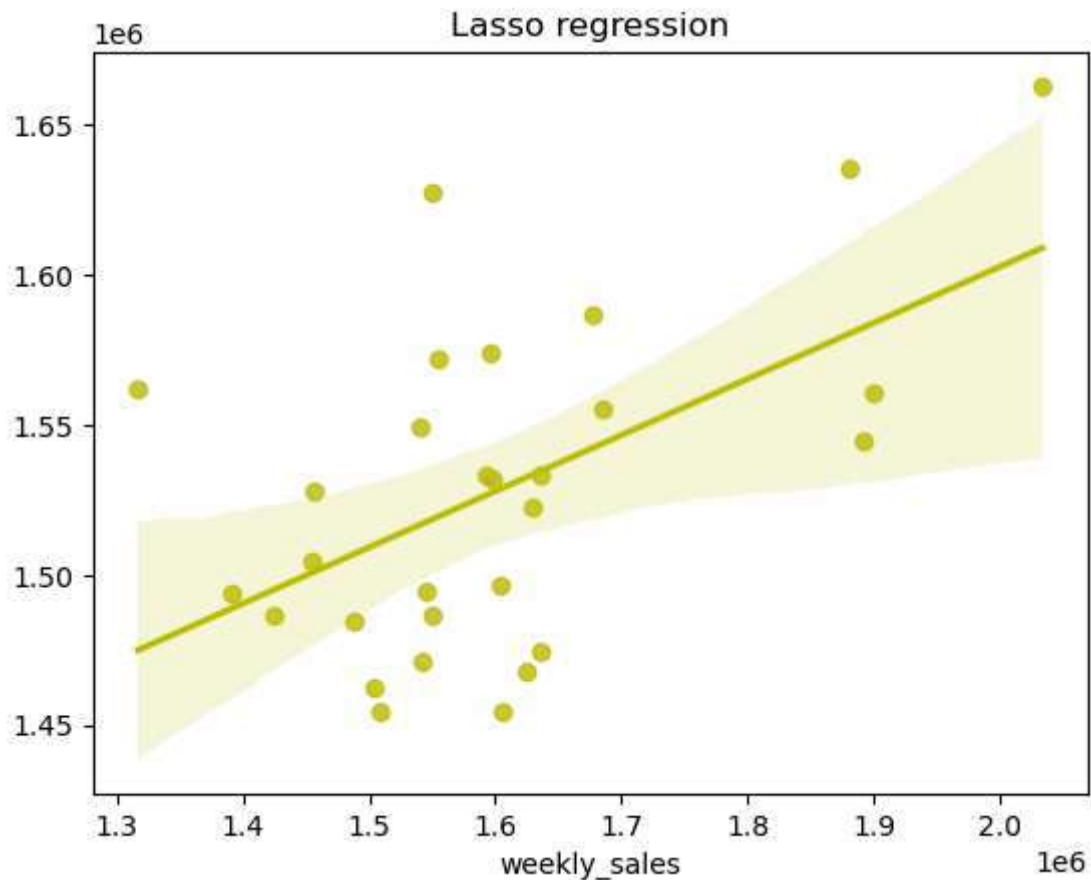
Root Mean Squared Error(RMSE):

152801.3455588893

R-Squared:

0.03468544187573952

```
Out[20]: <Axes: title={'center': 'Lasso regression'}, xlabel='weekly_sales'>
```



```
In [21]: ridge_reg= Ridge(alpha=1)
ridge_reg.fit(X_train,y_train)
ridge_y_pred= ridge_reg.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, ridge_y_pred)))
print('R-Squared:\n',r2(y_test,ridge_y_pred))
plt.title('Ridge regression')
sns.regplot(x=y_test,y=ridge_y_pred,color='y')
```


Root Mean Squared Error(RMSE):

153500.45817498674

R-Squared:

0.025832019326815225

```
Out[21]: <Axes: title={'center': 'Ridge regression'}, xlabel='weekly_sales'>
```



Comparing Coefficient of Regression

In [22]: `model.coef_`

Out[22]: `array([61993.93468347, -2617.60358042, -34338.15293787, 13422.2576375 ,
42279.48782025])`

In [23]: `lasso_reg.coef_`

Out[23]: `array([61981.19739554, -2617.81139176, -34315.03465077, 13418.36841804,
42250.22731354])`

In [24]: `ridge_reg.coef_`

Out[24]: `array([55418.46643272, -2685.2272262 , -28648.03740797, 12409.08802479,
34293.58901002])`

Regression Coefficient for Fuel_Price is reducing accuracy of model in comparison with CPI and Unemployment so we drop it

In [25]: `y= data.weekly_sales
X= data.drop(columns=['weekly_sales','store','date','fuel_price'])
X_train,X_test,y_train,y_test= tts(X,y,test_size=0.2,random_state=42)`

In [26]: `model= LinearRegression()
model.fit(X_train,y_train)
y_test_pred= model.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, y_test_pred)))
print('R-Squared:\n',r2(y_test,y_test_pred))`

```
plt.title('Linear regression')
sns.regplot(x=y_test,y=y_test_pred,color='green')
```


Root Mean Squared Error(RMSE):

152260.04079652095

R-Squared:

0.041512656979042495

Out[26]: <Axes: title={'center': 'Linear regression'}, xlabel='weekly_sales'>



```
In [27]: lasso_reg=Lasso(alpha=1)
lasso_reg.fit(X_train,y_train)
lasso_y_pred=lasso_reg.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, lasso_y_pred)))
print('R-Squared:\n',r2(y_test,lasso_y_pred))
plt.title('Lasso regression')
sns.regplot(x=y_test,y=lasso_y_pred,color='y')
```


Root Mean Squared Error(RMSE):

152262.39307026667

R-Squared:

0.04148304130290403

Out[27]: <Axes: title={'center': 'Lasso regression'}, xlabel='weekly_sales'>



```
In [28]: ridge_reg= Ridge(alpha=1)
ridge_reg.fit(X_train,y_train)
ridge_y_pred= ridge_reg.predict(X_test)
print('Root Mean Squared Error(RMSE):\n', np.sqrt(mse(y_test, ridge_y_pred)))
print('R-Squared:\n',r2(y_test,ridge_y_pred))
plt.title('Ridge regression')
sns.regplot(x=y_test,y=ridge_y_pred,color='y')
```

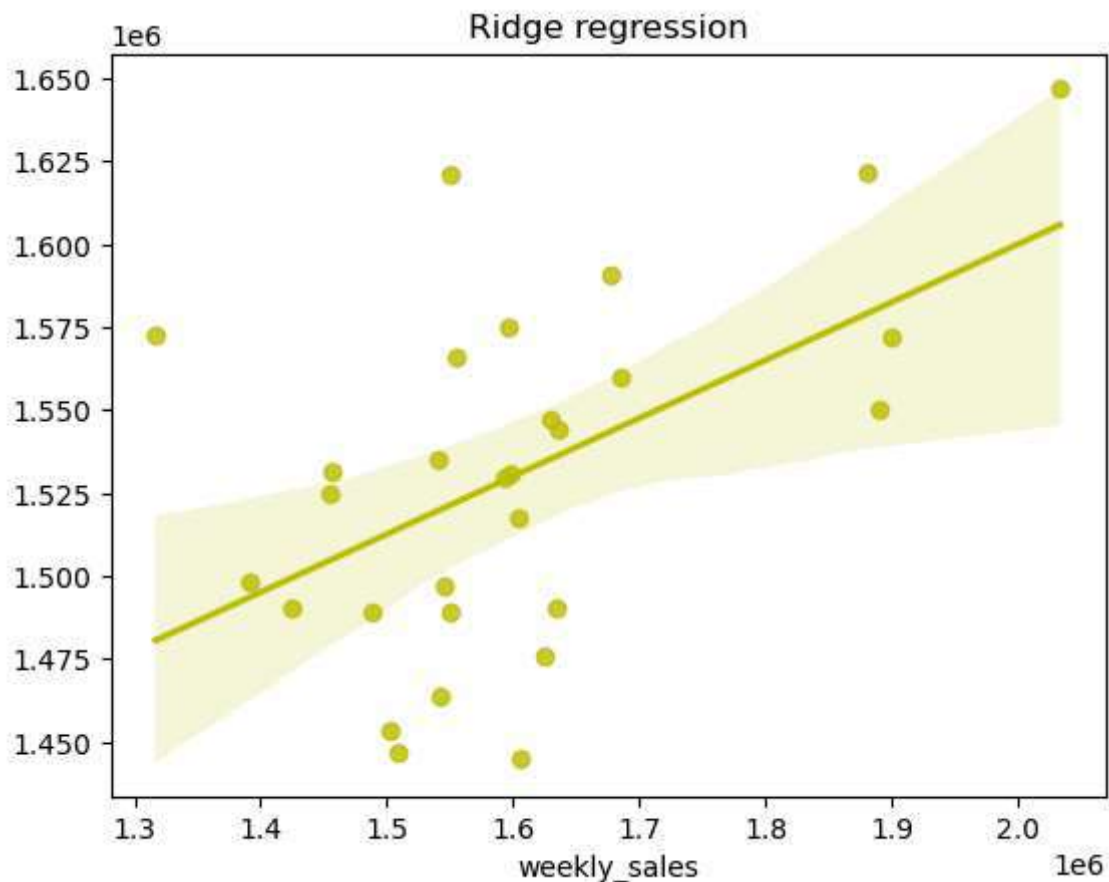
Root Mean Squared Error(RMSE):

153007.1958681327

R-Squared:

0.032082792684975825

```
Out[28]: <Axes: title={'center': 'Ridge regression'}, xlabel='weekly_sales'>
```



Linear Regression Model give the best performance on test case

2.2 Creating Day column from Date in the Dataframe

We cannot utilize DAY column in ML models as there is only 1 unique value in the whole column

```
In [29]: df['days'] = df.date.dt.dayofweek
df.days = df.days.map({0: 'Monday',
                        1: 'Tuesday',
                        2: 'Wednesday',
                        3: 'Thursday',
                        4: 'Friday',
                        5: 'Saturday',
                        6: 'Friday',})
```

```
In [30]: df.head()
```

Out[30]:

	store	date	weekly_sales	holiday_flag	temp	fuel_price	cpi	unemployment	days
0	1	2010-02-05	1643690.90	0	42.31	2.572	211.096358	8.106	Friday
1	1	2010-02-12	1641957.44	1	38.51	2.548	211.242170	8.106	Friday
2	1	2010-02-19	1611968.17	0	39.93	2.514	211.289143	8.106	Friday
3	1	2010-02-26	1409727.59	0	46.63	2.561	211.319643	8.106	Friday
4	1	2010-03-05	1554806.68	0	46.50	2.625	211.350143	8.106	Friday

In []: